

16-720 Computer Vision: Homework 3 (Spring 2021)

Lucas-Kanade Tracking

Instructor: Deva Ramanan

TAs: Michael Lee, Zhiqiu Lin, Cormac O'Meara, Thomas Weng, Gengshan Yang, Jason Zhang

This homework consists of three sections. In the first section you will implement a simple Lucas-Kanade (LK) tracker with one single template. The second section requires you to implement a motion subtraction method for tracking moving pixels in a scene. In the final section you shall study efficient tracking such as inverse composition. Other than the course slide decks, the following references may also be helpful:

1. Simon Baker, et al. *Lucas-Kanade 20 Years On: A Unifying Framework: Part 1*, CMU-RI-TR-02-16, Robotics Institute, Carnegie Mellon University, 2002
2. Simon Baker, et al. *Lucas-Kanade 20 Years On: A Unifying Framework: Part 2*, CMU-RI-TR-03-35, Robotics Institute, Carnegie Mellon University, 2003

Both are available at:

https://www.ri.cmu.edu/pub_files/pub3/baker_simon_2002_3/baker_simon_2002_3.pdf.

<https://www.ri.cmu.edu/publications/lucas-kanade-20-years-on-a-unifying-framework-part-2/>.

Instructions

1. **Integrity and collaboration:** Students are encouraged to work in groups but each student must submit their own work. If you work as a group, include the names of your collaborators in your write-up. Code should not be shared or copied. Do not use external code unless permitted (if in doubt, use Piazza to clarify whether a package can be used). Plagiarism is strongly prohibited and may lead to failure of this course.
2. **Start early!**
3. **Questions:** If you have any questions, please look at Piazza first. Other students may have encountered the same problem, and it may be solved already. If not, post your question on the discussion board. Please keep questions about homework problems in the corresponding threads.
4. The assignment must be completed using Python 3. We recommend setting up a conda environment. See this Piazza post for instructions. You will need the cv2, and scipy modules. You can install these packages with `pip install opencv-python scipy`.
5. Please stick to the function prototypes mentioned in the handout. This makes verifying code easier for the TAs.
6. **File paths:** Please make sure that any file paths that you use are relative and not absolute. Not `cv2.imread('/name/Documents/subdirectory/hw3/data/xyz.jpg')` but `cv2.imread('../data/xyz.jpg')`.

7. **Write-up:** Your write-up should mainly consist of four parts: your answers to theory questions, resulting images of each step, the discussions for experiments, and **as well as snippets of all your code in an Appendix** (*make sure that there aren't any lines that run beyond the page width and that all code is visible, else you may be penalized up to 5% of the total score*). Please note that we will not accept handwritten scans for your write-up. You may complete your write-up using L^AT_EX, Microsoft Word, or in-line on a Jupyter Notebook. We will provide a Latex template, but you are not required to use it as long as your submission follows all of the submission guidelines.
8. In your PDF, *add a page break after each question*. **When submitting to Gradescope, make sure that you select each page corresponding to your answer for each question (if the corresponding write-up and code live on separate pages, select all of them).** Not doing this makes it difficult for us to find your answer and you will be penalized accordingly.
9. **Submission:** Create a zip file, `<andrew-id>.zip`, composed of your write-up, your Python implementations (including helper functions) and results, and the implementations and results for extra credits (optional). Please make sure to remove the `data/` folder, and any other temporary files you've generated. Your final upload should have the files arranged in this layout:

- `<AndrewId>.zip`
 - `<AndrewId>/`
 - * `<AndrewId>.pdf`
 - * `code/`
 - `LucasKanade.py`
 - `LucasKanadeAffine.py`
 - `SubtractDominantMotion.py`
 - `InverseCompositionAffine.py`
 - `testCarSequence.py`
 - `testCarSequenceWithTemplateCorrection.py`
 - `testGirlSequence.py`
 - `testGirlSequenceWithTemplateCorrection.py`
 - `testAntSequence.py`
 - `testAerialSequence.py`
 - `LucasKanadePyramid.py` (*optional for extra credit*)
 - `testCarSequencePyramid.py` (*optional for extra credit*)
 - * `results/`
 - `carseqrects.npy`
 - `carseqrects-wcrt.npy`
 - `girlseqrects.npy`
 - `girlseqrects-wcrt.npy`

Please make sure you follow the submission rules mentioned above before uploading your zip file to Canvas. Assignments that violate this submission rule will be penalized by up to 5% of the total score.

1 Lucas-Kanade Tracking

In this section you will be implementing a simple Lucas & Kanade tracker with one single template. In the scenario of two-dimensional tracking with a pure translation warp function,

$$\mathcal{W}(\mathbf{x}; \mathbf{p}) = \mathbf{x} + \mathbf{p} . \quad (1)$$

The problem can be described as follows: starting with a rectangular neighborhood of pixels $\mathbb{N} \in \{\mathbf{x}_d\}_{d=1}^D$ on frame \mathcal{I}_t , the Lucas-Kanade tracker aims to move it by an offset $\mathbf{p} = [p_x, p_y]^T$ to obtain another rectangle on frame \mathcal{I}_{t+1} , so that the pixel squared difference in the two rectangles is minimized:

$$\mathbf{p}^* = \arg \min_{\mathbf{p}} \sum_{\mathbf{x} \in \mathbb{N}} \|\mathcal{I}_{t+1}(\mathbf{x} + \mathbf{p}) - \mathcal{I}_t(\mathbf{x})\|_2^2 \quad (2)$$

$$= \left\| \begin{bmatrix} \mathcal{I}_{t+1}(\mathbf{x}_1 + \mathbf{p}) \\ \vdots \\ \mathcal{I}_{t+1}(\mathbf{x}_D + \mathbf{p}) \end{bmatrix} - \begin{bmatrix} \mathcal{I}_t(\mathbf{x}_1) \\ \vdots \\ \mathcal{I}_t(\mathbf{x}_D) \end{bmatrix} \right\|_2^2 \quad (3)$$

Q1.1 (5 points) Starting with an initial guess of \mathbf{p} (for instance, $\mathbf{p} = [0, 0]^T$), we can compute the optimal \mathbf{p}^* iteratively. In each iteration, the objective function is locally linearized by first-order Taylor expansion,

$$\mathcal{I}_{t+1}(\mathbf{x}' + \Delta \mathbf{p}) \approx \mathcal{I}_{t+1}(\mathbf{x}') + \frac{\partial \mathcal{I}_{t+1}(\mathbf{x}')}{\partial \mathbf{x}'^T} \frac{\partial \mathcal{W}(\mathbf{x}; \mathbf{p})}{\partial \mathbf{p}^T} \Delta \mathbf{p} \quad (4)$$

where $\Delta \mathbf{p} = [\Delta p_x, \Delta p_y]^T$, is the template offset. Further, $\mathbf{x}' = \mathcal{W}(\mathbf{x}; \mathbf{p}) = \mathbf{x} + \mathbf{p}$ and $\frac{\partial \mathcal{I}(\mathbf{x}')}{\partial \mathbf{x}'^T}$ is a vector of the x - and y - image gradients at pixel coordinate \mathbf{x}' . In a similar manner to Equation 3 one can incorporate these linearized approximations into a vectorized form such that,

$$\arg \min_{\Delta \mathbf{p}} \|\mathbf{A} \Delta \mathbf{p} - \mathbf{b}\|_2^2 \quad (5)$$

such that $\mathbf{p} \leftarrow \mathbf{p} + \Delta \mathbf{p}$ at each iteration.

- What is $\frac{\partial \mathcal{W}(\mathbf{x}; \mathbf{p})}{\partial \mathbf{p}^T}$?
- What is \mathbf{A} and \mathbf{b} ?
- What conditions must $\mathbf{A}^T \mathbf{A}$ meet so that a unique solution to $\Delta \mathbf{p}$ can be found?

Q1.2 (15 points) Implement a function with the following signature

```
LucasKanade(It, It1, rect, p0 = np.zeros(2))
```

that computes the optimal local motion from frame \mathcal{I}_t to frame \mathcal{I}_{t+1} that minimizes Equation 3. Here \mathbf{It} is the image frame \mathcal{I}_t , $\mathbf{It1}$ is the image frame \mathcal{I}_{t+1} , \mathbf{rect} is the 4-by-1 vector that represents a rectangle describing all the pixel coordinates within \mathbb{N} within the image frame \mathcal{I}_t , and p_0 is the initial parameter guess $(\delta x, \delta y)$. The four components of the rectangle are $[x_1, y_1, x_2, y_2]^T$, where $[x_1, y_1]^T$ is the top-left corner and $[x_2, y_2]^T$ is the bottom-right corner. The rectangle is inclusive, i.e., it includes all the four corners. To deal with fractional movement of the template, you will need to interpolate the image using the Scipy module `ndimage.shift` or something similar. You will also need to iterate the

estimation until the change in $\|\Delta \mathbf{p}\|_2^2$ is below a threshold. In order to perform interpolation you might find `RectBivariateSpline` from the `scipy.interpolate` package. Read the documentation of defining the spline (`RectBivariateSpline`) as well as evaluating the spline using `RectBivariateSpline.ev` carefully.

Q1.3 (10 points) Write a script `testCarSequence.py` that loads the video frames from `carseq.npy`, and runs the Lucas-Kanade tracker that you have implemented in the previous task to track the car. `carseq.npy` can be located in the `data` directory and it contains one single three-dimensional matrix: the first two dimensions correspond to the *height* and *width* of the frames respectively, and the third dimension contain the indices of the frames (that is, the first frame can be visualized with `imshow(frames[:, :, 0])`). The rectangle in the first frame is $[x_1, y_1, x_2, y_2]^T = [59, 116, 145, 151]^T$.

- Report your tracking performance (image + bounding rectangle) at frames 1, 100, 200, 300 and 400 in a format similar to Figure 1.
- Create a file called `carseqrects.npy`, which contains one single $n \times 4$ matrix, where each row stores the `rect` that you have obtained for each frame, and n is the total number of frames. You are encouraged to play with the parameters defined in the scripts and report the best results.

Similarly, write a script `testGrilSequence.py` that loads the video from `girlseq.npy` and runs your Lucas-Kanade tracker on it. The rectangle in the first frame is $[x_1, y_1, x_2, y_2]^T = [280, 152, 330, 318]^T$.

- Report your tracking performance (image + bounding rectangle) at frames 1, 20, 40, 60 and 80 in a format similar to Figure 1.
- Also, create a file called `girlseqrects.npy`, which contains one single $n \times 4$ matrix, where each row stores the `rect` that you have obtained for each frame.




Figure 1: Lucas-Kanade Tracking with One Single Template

Q1.4 (20 points) As you might have noticed, the image content we are tracking in the first frame differs from the one in the last frame. This is understandable since we are updating the template after processing each frame and the error can be accumulating. This problem is known as *template drifting*. There are several ways to mitigate this problem. Iain Matthews et al. (2003, https://www.ri.cmu.edu/publication_view.html?pub_id=4433) suggested one possible approach. Write two scripts with a similar functionality to **Q1.3** but with a template correction routine incorporated: `testCarSequenceWithTemplateCorrection.py` and

`testGirlSequenceWithTemplateCorrection.py`. Save the rects as `carseqrects-wcrt.npy` and `girlseqrects-wcrt.npy`, and also report the performance at those frames. An example is given in Figure 2. Again, you are encouraged to play with the parameters defined in the scripts to see how each parameter affects the tracking results.




Figure 2: Lucas-Kanade Tracking with Template Correction

Here the blue rectangles are created with the baseline tracker in **Q1.3**, the red ones with the tracker in **Q1.4**. The tracking performance has been improved non-trivially. Note that you do not necessarily have to draw two rectangles in each frame, but make sure that the performance improvement can be easily visually inspected.

2 Affine Motion Subtraction

In this section, you will implement a tracker for estimating dominant affine motion in a sequence of images and subsequently identify pixels corresponding to moving objects in the scene. You will be using the images in the file `aerialseq.npy`, which consists of aerial views of moving vehicles from a non-stationary camera.

2.1 Dominant Motion Estimation

In the first section of this homework we assumed the the motion is limited to pure translation. In this section you shall implement a tracker for affine motion using a planar affine warp function. To estimate dominant motion, the entire image \mathcal{I}_t will serve as the template to be tracked in image \mathcal{I}_{t+1} , that is, \mathcal{I}_{t+1} is assumed to be approximately an affine warped version of \mathcal{I}_t . This approach is reasonable under the assumption that a majority of the pixels correspond to the stationary objects in the scene whose depth variation is small relative to their distance from the camera.

Using a planar affine warp function you can recover the vector $\Delta \mathbf{p} = [p_1, \dots, p_6]^T$,

$$\mathbf{x}' = \mathcal{W}(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} 1 + p_1 & p_2 \\ p_4 & 1 + p_5 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} p_3 \\ p_6 \end{bmatrix}. \quad (6)$$

One can represent this affine warp in homogeneous coordinates as,

$$\tilde{\mathbf{x}}' = \mathbf{M} \tilde{\mathbf{x}} \quad (7)$$

where,

$$\mathbf{M} = \begin{bmatrix} 1 + p_1 & p_2 & p_3 \\ p_4 & 1 + p_5 & p_6 \\ 0 & 0 & 1 \end{bmatrix}. \quad (8)$$

Here \mathbf{M} represents $\mathbf{W}(\mathbf{x}; \mathbf{p})$ in homogeneous coordinates as described in [1]. Also note that \mathbf{M} will differ between successive image pairs. Starting with an initial guess of $\mathbf{p} = \mathbf{0}$ (i.e. $\mathbf{M} = \mathbf{I}$) you will need to solve a sequence of least-squares problem to determine $\Delta \mathbf{p}$ such that $\mathbf{p} \rightarrow \mathbf{p} + \Delta \mathbf{p}$ at each iteration. Note that unlike previous examples where the template to be tracked is usually small in comparison with the size of the image, image \mathcal{I}_t will almost always not be contained fully in the warped version \mathcal{I}_{t+1} . Hence, one must only consider pixels lying in the region common to \mathcal{I}_t and the warped version of \mathcal{I}_{t+1} when forming the linear system at each iteration.

Q2.1 (15 points) Write a function with the following signature

`LucasKanadeAffine(It, It1)`

which returns the affine transformation matrix \mathbf{M} , and It and It1 are \mathcal{I}_t and \mathcal{I}_{t+1} respectively. `LucasKanadeAffine` should be relatively similar to `LucasKanade` from the first section (you will probably also find `scipy.ndimage.affine_transform` helpful).

2.2 Moving Object Detection

Once you are able to compute the transformation matrix \mathbf{M} relating an image pair \mathcal{I}_t and \mathcal{I}_{t+1} , a naive way for determining pixels lying on moving objects is as follows: warp the image \mathcal{I}_t using \mathbf{M} so that it is registered to \mathcal{I}_{t+1} and subtract it from \mathcal{I}_{t+1} ; the locations where the absolute difference exceeds a threshold can then be declared as corresponding to locations of moving objects. To obtain better results, you can check out the following `scipy.morphology` functions: `binary_erosion`, and `binary_dilation`.




Figure 3: Lucas-Kanade Tracking with Motion Detection

Q2.2 (10 points) Using the function you have developed for dominant motion estimation, write a function with the following signature

`SubtractDominantMotion(image1, image2)`

where `image1` and `image2` form the input image pair, and the return value `mask` is a binary image of the same size that dictates which pixels are considered to be corresponding to moving objects. You should invoke `LucasKanadeAffine` in this function to derive the transformation matrix \mathbf{M} , and produce the aforementioned binary mask accordingly.

Q2.3 (10 points) Write two scripts `testAntSequence.py` and `testAerialSequence.py` that load the image sequence from `antseq.npy` and `aerialseq.npy` and run the motion detection routine you have developed to detect the moving objects. Try to implement `testAntSequence.py` first as it involves little camera movement and can help you debug your mask generation procedure. **Report the performance at frames 30, 60, 90 and 120 with the corresponding binary masks superimposed, as exemplified in Figure 3.** Feel free to visualize the motion detection performance in a way that you would prefer, but please make sure it can be visually inspected without undue effort.

3 Efficient Tracking

3.1 Inverse Composition

The inverse compositional extension of the Lucas-Kanade algorithm (see [1]) has been used in literature to great effect for the task of efficient tracking. When utilized within tracking it attempts to linearize the current frame as,

$$\mathcal{I}_t(\mathcal{W}(\mathbf{x}; \mathbf{0} + \Delta\mathbf{p}) \approx \mathcal{I}_t(\mathbf{x}) + \frac{\partial \mathcal{I}_t(\mathbf{x})}{\partial \mathbf{x}^T} \frac{\partial \mathcal{W}(\mathbf{x}; \mathbf{0})}{\partial \mathbf{p}^T} \Delta\mathbf{p} . \quad (9)$$

In a similar manner to the conventional Lucas-Kanade algorithm one can incorporate these linearized approximations into a vectorized form such that,

$$\arg \min_{\Delta\mathbf{p}} \|\mathbf{A}'\Delta\mathbf{p} - \mathbf{b}'\|_2^2 \quad (10)$$

for the specific case of an affine warp where $\mathbf{p} \leftarrow \mathbf{M}$ and $\Delta\mathbf{p} \leftarrow \Delta\mathbf{M}$ this results in the update $\mathbf{M} = \mathbf{M}(\Delta\mathbf{M})^{-1}$. Just to clarify, the notation $\mathbf{M}(\Delta\mathbf{M})^{-1}$ corresponds to $\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p})^{-1}; \mathbf{p})$ in Section 2.2 from [2].

Q3.1 (10 points) Reimplement the function `LucasKanadeAffine(It, It1)` as `InverseCompositionAffine(It, It1)` using the inverse compositional method.

Q3.2 (5 points) In your own words please describe why the inverse compositional approach is more computationally efficient than the classical approach?

4 Extra Credit: LK Tracking on an Image Pyramid

If the target being tracked moves a lot between frames, the LK tracker can break down. One way to mitigate this problem is to run the LK tracker on a set of image pyramids instead of a single image. The Pyramid tracker starts by performing tracking on a higher level (smaller image) to get a coarse alignment estimate, propagating this down into the lower level and repeating until a fine aligning warp has been found at the lowest level of the pyramid (the original sized image). In addition to being more robust, the pyramid version of the tracker is much faster because it needs to run fewer gradient descent iterations on the full scale image due to its coarse to fine approach.

Q4.1 (10 points) Implement these modifications for the Lucas-Kanade tracker that you implemented for Q1.2. Use the same function names with Pyramid appended. Compare the performance of your pyramid algorithm to the original tracker on the car sequence under 3 conditions:

- Skipping every 2nd frame. Plot iterations [1, 50, 100, 150, 200].
- Skipping every 3rd frame. Plot iterations [1, 35, 70, 105, 140].
- Skipping every 4th frame. Plot iterations [1, 25, 50, 75, 100].

Compare the two algorithms in terms of computational performance and tracking accuracy.

5 Frequently Asked Questions (FAQs)

Q1: Why do we need to use `ndimage.shift` or `RectBivariateSpline` for moving the rectangle template?

A1: When moving the rectangle template with $\Delta \mathbf{p}$, you can either move the points inside the template or move the image in the opposite direction. If you choose to move the points, the new points can have fractional coordinates, so you need to use `RectBivariateSpline` to sample the image intensity at those fractional coordinates. If you instead choose to move the image with `ndimage.shift`, you don't need to move the points and you can sample the image intensity at those points directly. The first approach could be faster since it does not require moving the entire image.

Q2: What's the right way of computing the image gradients $\mathcal{I}_x(\mathbf{x})$ and $\mathcal{I}_y(\mathbf{x})$. Should I first sample the image intensities $\mathcal{I}(\mathbf{x})$ at \mathbf{x} and then compute the image gradients $\mathcal{I}_x(\mathbf{x})$ and $\mathcal{I}_y(\mathbf{x})$ with $\mathcal{I}(\mathbf{x})$? Or should I first compute the entire image gradients \mathcal{I}_x and \mathcal{I}_y and sample them at \mathbf{x} ?

A2: The second approach is the correct one.

Q3: Can I use pseudo-inverse the least-squared problem $\arg \min_{\Delta \mathbf{p}} \|\mathbf{A} \Delta \mathbf{p} - \mathbf{b}\|_2^2$?

A3: Yes, the pseudo-inverse solution of $\mathbf{A} \Delta \mathbf{p} = \mathbf{b}$ is also $\Delta \mathbf{p} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}$ when \mathbf{A} has full column ranks, i.e., the linear system is overdetermined.

Q4: For inverse compositional Lucas Kanade, how to deal with points outside out the image?

A4: Since the Hessian in inverse compositional Lucas Kanade is precomputed, we cannot simply remove points when they are outside the image since it can result in dimension mismatch. However, we can set the error $\mathcal{I}_{t+1}(\mathcal{W}(\mathbf{x}; \mathbf{p})) - \mathcal{I}_t(\mathbf{x})$ to 0 for \mathbf{x} outside the image.