16-720 Computer Vision: Homework 4 (Spring 2021)
3D Reconstruction

Instructor: Deva Ramanan

TAs: Michael Lee, Zhiqiu Lin, Cormac O'Meadhra, Thomas Weng, Gengshan Yang, Jason Zhang

Instructions

1. **Integrity and collaboration:** Students are encouraged to work in groups but each student must submit their own work. If you work as a group, include the names of your collaborators in your write-up. Code should not be shared or copied. Do not use external code unless permitted (if in doubt, use Piazza to clarify whether a package can be used). Plagiarism is strongly prohibited and may lead to failure of this course.

2. **Start early!**

3. **Questions:** If you have any questions, please look at Piazza first. Other students may have encountered the same problem, and it may be solved already. If not, post your question on the discussion board. Please keep questions about homework problems in the corresponding threads.

4. The assignment must be completed using Python 3. We recommend setting up a conda environment. See this Piazza post for instructions. You will need the cv2, and scipy modules. You can install these packages with `pip install opencv-python scipy`.

5. Please stick to the function prototypes mentioned in the handout. This makes verifying code easier for the TAs.

6. **File paths:** Please make sure that any file paths that you use are relative and not absolute. Not `cv2.imread('/name/Documents/subdirectory/hw/data/xyz.jpg')` but `cv2.imread('../data/xyz.jpg')`.

7. **Write-up:** Your write-up should mainly consist of four parts: your answers to theory questions, resulting images of each step, the discussions for experiments, and **as well as snippets of all your code in an Appendix** *(make sure that there aren't any lines that run beyond the page width and that all code is visible, else you may be penalized).* Please note that we will not accept handwritten scans for your write-up. You may complete your write-up using LaTeX, Microsoft Word, or in-line on a Jupyter Notebook. We will provide a Latex template, but you are not required to use it as long as your submission follows all of the submission guidelines.

8. In your PDF, *add a page break after each question.* **When submitting to Gradescope, make sure that you select each page corresponding to your answer for each question (*if the corresponding write-up and code live on separate pages, select all of them*).** Not doing this makes it difficult for us to find your answer and you will be penalized accordingly.

9. **Submission:** Create a zip file, `<andrew-id>.zip`, composed of your write-up, your Python implementations (including helper functions) and results. Please make sure to remove the `data/` folder, and any other temporary files you've generated. Your final upload should have the files arranged in this layout:

- <ins>`<AndrewId>.zip`</ins>
  - `<AndrewId>/`
    * `<AndrewId>.pdf`
    * <ins>`code/`</ins>
      · `submission.py`: your implementation of algorithms.
      · `findM2.py`: script to compute the correct camera matrix.
      · `visualize.py`: script to visualize the 3d points.
    * <ins>`results/`</ins>
      · `q2_1.npz`: file with output of Q2.1.
      · `q3_3.npz`: file with output of Q3.3.
      · `q4_1.npz`: file with output of Q4.1.
      · `q4_2.npz`: file with output of Q4.2.

Please follow the submission rules mentioned above before uploading your zip file to Gradescope. Not doing this will incur a significant penalty.

Part I

Theory

Before implementing our own 3D reconstruction, let's take a look at some simple theory questions that may arise. The answers to the below questions should be relatively short, consisting of a few lines of math and text (maybe a diagram if it helps your understanding).

Q1.1 **(5 points)** Suppose two cameras fixate on a point $\mathbf{x}$ (see Figure 1) in space such that their principal axes intersect at that point. Show that if the image coordinates are normalized so that the coordinate origin $(0, 0)$ coincides with the principal point, the $\mathbf{F}_{33}$ element of the fundamental matrix is zero.


Figure 1: Figure for Q1.1. $C1$ and $C2$ are the optical centers. The principal axes intersect at point $\mathbf{w}$ (P in the figure).

Q1.2 **(5 points)** Consider the case of two cameras viewing an object such that the second camera differs from the first by a *pure translation* that is parallel to the x-axis. Show that the epipolar lines in the two cameras are also parallel to the x-axis. Backup your argument with relevant equations.

Q1.3 **(5 points)** Suppose we have an inertial sensor which gives us the accurate positions ($\mathbf{R}_i$ and $\mathbf{t}_i$, the rotation matrix and translation vector) of the robot at time τ_i. What will be the effective rotation ($\mathbf{R}_{rel}$) and translation ($\mathbf{t}_{rel}$) between two frames at different time stamps, τ_i, τ_j? Suppose the camera intrinsics ($\mathbf{K}$) are known, express the essential matrix ($\mathbf{E}$) and the fundamental matrix ($\mathbf{F}$) in terms of $\mathbf{K}$, $\mathbf{R}_{rel}$ and $\mathbf{t}_{rel}$.

Q1.4 **(extra credit, 10 points)** Suppose that a camera views an object and its reflection in a plane mirror. Show that this situation is equivalent to having two images of the object which are related by a skew-symmetric fundamental matrix. You may assume that the object is flat, meaning that all points on the object are of equal distance to the mirror (*Hint:* draw the relevant vectors to understand the relationships between the camera, the object and its reflected image.)

Figure 2: Figure for Q1.3. $C1$ and $C2$ are the optical centers. The rotation and the translation is obtained using inertial sensors. $\mathbf{R}_{rel}$ and $\mathbf{t}_{rel}$ are the relative rotation and translation between two frames.

Part II
Practice

1   Overview

In this part you will begin by implementing the two different methods seen in class to estimate the fundamental matrix from corresponding points in two images (Section 2). Next, given the fundamental matrix and calibrated intrinsics (which will be provided) you will compute the essential matrix and use this to compute a 3D metric reconstruction from 2D correspondences using triangulation (Section 3). Then, you will implement a method to automatically match points taking advantage of epipolar constraints and make a 3D visualization of the results (Section 4). Finally, you will implement RANSAC and bundle adjustment to further improve your algorithm (Section 5).

2   Fundamental matrix estimation

In this section you will explore different methods of estimating the fundamental matrix given a pair of images. In the `data/` directory, you will find two images (see Figure 3) from the Middlebury multi-view dataset[1], which is used to evaluate the performance of modern 3D reconstruction algorithms.

---

[1] `http://vision.middlebury.edu/mview/data/`

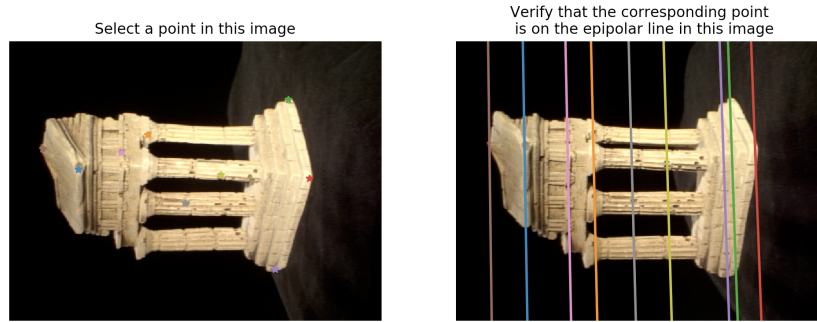Figure 3: `Temple` images for this assignment

Figure 4: `displayEpipolarF in helper.py` creates a GUI for visualizing epipolar lines

The Eight Point Algorithm

The 8-point algorithm (discussed in class, and outlined in Section 10.1 of Forsyth & Ponce) is arguably the simplest method for estimating the fundamental matrix. For this section, you can use provided correspondences you can find in `data/some_corresp.npz`.

Q2.1 **(10 points)** Finish the function `eightpoint` in `submission.py`. Make sure you follow the signature for this portion of the assignment:

$$F = eightpoint(pts1, pts2, M)$$

where `pts1` and `pts2` are $N \times 2$ matrices corresponding to the (x, y) coordinates of the N points in the first and second image repectively. `M` is a scale parameter.

- You should scale the data as was discussed in class, by dividing each coordinate by M (the maximum of the image's width and height). After computing $\mathbf{F}$, you will have to "unscale" the fundamental matrix.
  Hint: If $\mathbf{x}_{normalized} = \mathbf{T}\mathbf{x}$, then $\mathbf{F}_{unnormalized} = \mathbf{T}^T \mathbf{F} \mathbf{T}$.
  You must enforce the singularity condition of the $\mathbf{F}$ before unscaling.

- You may find it helpful to refine the solution by using local minimization. This probably won't fix a completely broken solution, but may make a good solution better by locally minimizing a geometric cost function.
  For this we have provided a helper function `refineF` in `util.py` taking in $\mathbf{F}$ and the two sets of points, which you can call from `eightpoint` before unscaling F.

- Remember that the x-coordinate of a point in the image is its column entry, and y-coordinate is the row entry. Also note that eight-point is just a figurative name, it just means that you need at least 8 points; your algorithm should use an over-determined system ($N > 8$ points).

- To visualize the correctness of your estimated $\mathbf{F}$, use the supplied function `displayEpipolarF` in `helper.py`, which takes in F, and the two images. This GUI lets you select a point in one of the images and visualize the corresponding epipolar line in the other image (Figure 4).

- **Output:** Save your matrix $\mathbf{F}$, scale `M` to the file `q2_1.npz`.
  In your write-up: Write your recovered $\mathbf{F}$ and include an image of some example output of `displayEpipolarF`.

3 Metric Reconstruction

You will compute the camera matrices and triangulate the 2D points to obtain the 3D scene structure. To obtain the Euclidean scene structure, first convert the fundamental matrix $\mathbf{F}$ to an essential matrix $\mathbf{E}$. Examine the lecture notes and the textbook to find out how to do this when the internal camera calibration matrices $\mathbf{K}_1$ and $\mathbf{K}_2$ are known; these are provided in `data/intrinsics.npz`.

Q3.1 **(5 points)** Write a function to compute the essential matrix $\mathbf{E}$ given $\mathbf{F}$, $\mathbf{K}_1$ and $\mathbf{K}_2$ with the signature:

$$E = \texttt{essentialMatrix(F, K1, K2)}$$

In your write-up: Write your estimated $\mathbf{E}$ using $\mathbf{F}$ from the eight-point algorithm.

Given an essential matrix, it is possible to retrieve the projective camera matrices $\mathbf{M}_1$ and $\mathbf{M}_2$ from it. Assuming $\mathbf{M}_1$ is fixed at $[\mathbf{I}, 0]$, $\mathbf{M}_2$ can be retrieved up to a scale and four-fold rotation ambiguity. For details on recovering $\mathbf{M}_2$, see section 7.2 in Szeliski. We have provided you with the function `camera2` in `code/helper.py` to recover the four possible $\mathbf{M}_2$ matrices given $\mathbf{E}$.
Note: The $\mathbf{M}1$ and $\mathbf{M}2$ here are projection matrices of the form: $\mathbf{M}_1 = \begin{bmatrix} \mathbf{I}|0 \end{bmatrix}$ and $\mathbf{M}_2 = \begin{bmatrix} \mathbf{R}|\mathbf{t} \end{bmatrix}$.

Q3.2 **(5+5 points)** Using the above, write a function to triangulate a set of 2D coordinates in the image to a set of 3D points with the signature:

$$[w, err] = \texttt{triangulate(C1, pts1, C2, pts2)}$$

where `pts1` and `pts2` are the $N \times 2$ matrices with the 2D image coordinates and `w` is an $N \times 3$ matrix with the corresponding 3D points per row. `C1` and `C2` are the 3×4 camera matrices. Remember that you will need to multiply the given intrinsics matrices with your solution for the canonical camera matrices to obtain the final camera matrices. Various methods exist for triangulation - probably the most familiar for you is based on least squares (see Szeliski Chapter 7 if you want to learn about other methods):

For each point i, we want to solve for 3D coordinates $\mathbf{w}_i = \left[x_i, y_i, z_i\right]^T$, such that when they are projected back to the two images, they are close to the original 2D points. To project the 3D coordinates back to 2D images, we first write $\mathbf{w}_i$ in homogeneous coordinates, and compute $\mathbf{C}_1 \tilde{\mathbf{w}}_i$ and $\mathbf{C}_2 \tilde{\mathbf{w}}_i$ to obtain the 2D homogeneous coordinates projected to camera 1 and camera 2, respectively.

For each point i, we can write this problem in the following form:

$$\mathbf{A}_i \mathbf{w}_i = 0,$$

where $\mathbf{A}_i$ is a 4×4 matrix, and $\tilde{\mathbf{w}}_i$ is a 4×1 vector of the 3D coordinates in the homogeneous form. Then, you can obtain the homogeneous least-squares solution (discussed in class) to solve for each $\mathbf{w}_i$.

Q3.2.1 **(5 points)** In your write-up, write down the expression for the matrix $\mathbf{A}_i$.

Q3.2.2 **(5 points)** Once you have implemented triangulation, check the performance by looking at the reprojection error:

$$\texttt{err} = \sum_i \|\mathbf{x}_{1i}, \widehat{\mathbf{x}_{1i}}\|^2 + \|\mathbf{x}_{2i}, \widehat{\mathbf{x}_{2i}}\|^2$$

where $\widehat{\mathbf{x}_{1i}} = Proj(\mathbf{C}_1, \mathbf{w}_i)$ and $\widehat{\mathbf{x}_{2i}} = Proj(\mathbf{C}_2, \mathbf{w}_i)$.
In your write-up: print out the scalar reproejction error defined as above.
Note: $\mathbf{C}_1$ and $\mathbf{C}_2$ here are projection matrices of the form: $\mathbf{C}_1 = \mathbf{K}_1 \mathbf{M}_1 = \mathbf{K}_1 \left[\mathbf{I}|0\right]$ and $\mathbf{C}_2 = \mathbf{K}_2 \mathbf{M}_2 = \mathbf{K}_2 \left[\mathbf{R}|\mathbf{t}\right]$.

Q3.3 **(10 points)** Write a script `findM2.py` to obtain the correct M2 from M2s by testing the four solutions through triangulations. Use the correspondences from `data/some_cor-resp.npz`.
Output: Save the correct M2, the corresponding C2, and 3D points P to `q3_3.npz`.
In your write-up: Write your recovered M2.

4 3D Visualization

You will now create a 3D visualization of the temple images. By treating our two images as a stereo-pair, we can triangulate corresponding points in each image, and render their 3D locations.

Q4.1 **(15 points)** Implement a function with the signature:

$$[\texttt{x2, y2}] = \texttt{epipolarCorrespondence(im1, im2, F, x1, y1)}$$

This function takes in the x and y coordinates of a pixel on `im1` and your fundamental matrix $\mathbf{F}$, and returns the coordinates of the pixel on `im2` which correspond to the input point. The match is obtained by computing the similarity of a small window around the (x_1, y_1) coordinates in `im1` to various windows around possible matches in the `im2` and returning the closest.

Instead of searching for the matching point at every possible location in `im2`, we can use $\mathbf{F}$ and simply search over the set of pixels that lie along the epipolar line (recall that the

epipolar line passes through a single point in `im2` which corresponds to the point (x_1, y_1) in `im1`).

There are various possible ways to compute the window similarity. For this assignment, simple methods such as the Euclidean or Manhattan distances between the intensity of the pixels should suffice. See Szeliski Chapter 11, on stereo matching, for a brief overview of these and other methods.

Implementation hints:

- Experiment with various window sizes.

- It may help to use a Gaussian weighting of the window, so that the center has greater influence than the periphery.

- Since the two images only differ by a small amount, it might be beneficial to consider matches for which the distance from (x_1, y_1) to (x_2, y_2) is small.

To help you test your `epipolarCorrespondence`, we have included a helper function `epipolarMatchGUI` in `code/helper.py`, which takes in two images the fundamental matrix. This GUI allows you to click on a point in `im1`, and will use your function to display the corresponding point in `im2`. See Figure 5.
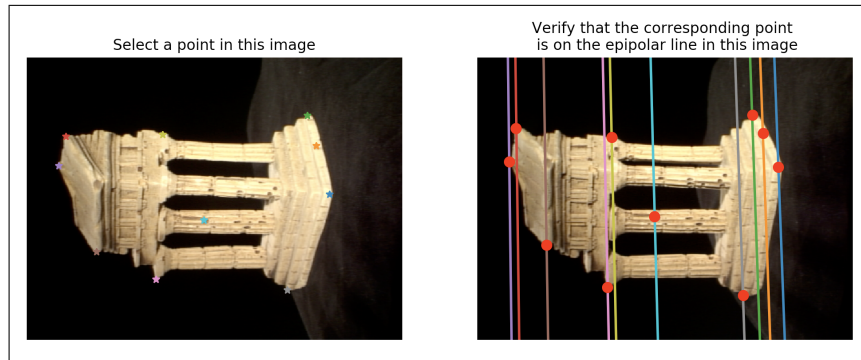

Figure 5: `epipolarMatchGUI` shows the corresponding point found by calling `epipolarCorrespondence`

It's not necessary for your matcher to get *every* possible point right, but it should get easy points (such as those with distinctive, corner-like windows). It should also be good enough to render an intelligible representation in the next question.

Output: Save the matrix **F**, points `pts1` and `pts2` which you used to generate the screenshot to the file `q4_1.npz`.

In your write-up: Include a screenshot of `epipolarMatchGUI` with some detected correspondences.

Q4.2 (10 points) Included in this homework is a file `data/templeCoords.npz` which contains 288 hand-selected points from `im1` saved in the variables `x1` and `y1`.

Now, we can determine the 3D location of these point correspondences using the `triangulate` function. These 3D point locations can then plotted using the Matplotlib or plotly package. Write a script `visualize.py`, which loads the necessary files from `../data/`
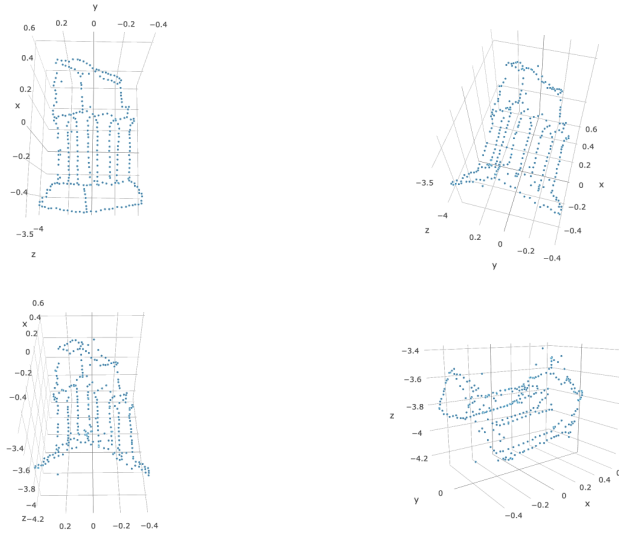
Figure 6: An example point cloud

to generate the 3D reconstruction using `scatter`. An example is shown in Figure 6.

Output: Again, save the matrix $\mathbf{F}$, matrices $\mathbf{M}1, \mathbf{M}2, \mathbf{C}1, \mathbf{C}2$ which you used to generate the screenshots to the file `q4_2.npz`.

In your write-up: Take a few screenshots of the 3D visualization so that the outline of the temple is clearly visible, and include them with your homework submission.

5 Bundle Adjustment (extra credit)

Q5.1 (15 points) In some real world applications, manually determining correspondences is infeasible and often there will be noisy coorespondances. Fortunately, the RANSAC method seen in class can be applied to the problem of fundamental matrix estimation.

Implement the above algorithm with the signature:

```
[F, inliers] = ransacF(pts1, pts2, M, nIters, tol)
```

where `M` is defined in the same way as in Section 2 and inliers is a boolean vector of size equivalent to the number of points. Here inliers is set to true only for the points that satisfy the threshold defined for the given fundamental matrix F.

We have provided some noisy coorespondances in `some_corresp_noisy.npz` in which around 75% of the points are inliers. Compare the result of RANSAC with the result of the eightpoint when ran on the noisy coorespondances. Briefly explain the error metrics you used, how you decided which points were inliers and any other optimizations you may have made. `nIters` is the maximum number of iterations of The RANSAC and `tol` is the tolerence of the error to be considered as inliers. Discuss the effect on the Fundamental matrix by varying these values.

- *Hints:* Use the Eight point to compute the fundamental matrix from the minimal set of points. Then compute the inliers, and refine your estimate using all the inliers.

9

Q5.2 (15 points) So far we have independently solved for camera matrix, $\mathbf{M}_j$ and 3D projections, $\mathbf{w}_i$. In bundle adjustment, we will jointly optimize the reprojection error with respect to the points $\mathbf{w}_i$ and the camera matrix $\mathbf{w}_j$.

$$err = \sum_{ij} \|\mathbf{x}_{ij} - Proj(\mathbf{C}_j, \mathbf{w}_i)\|^2,$$

where $\mathbf{w}_j = \mathbf{K}_j \mathbf{M}_j$, same as in Q3.2.

For this homework we are going to only look at optimizing the extrinsic matrix. To do this we will be parametrizing the rotation matrix $\mathbf{R}$ using Rodrigues formula to produce vector $\mathbf{r} \in \mathbb{R}^3$. Write a function that converts a Rodrigues vector $\mathbf{r}$ to a rotation matrix $\mathbf{R}$

$$R = \text{rodrigues(r)}$$

as well as the inverse function that converts a rotation matrix $\mathbf{R}$ to a Rodrigues vector $\mathbf{r}$

$$r = \text{invRodrigues(R)}.$$

Q5.3 (10 points) Using this parameterization, write an optimization function

```
residuals = rodriguesResidual(K1, M1, p1, K2, p2, x)
```

where x is the flattened concatenation of $\mathbf{x}$, $\mathbf{r}_2$, and $\mathbf{t}_2$. $\mathbf{w}$ are the 3D points; $\mathbf{r}_2$ and $\mathbf{t}_2$ are the rotation (in the Rodrigues vector form) and translation vectors associated with the projection matrix $\mathbf{M}_2$. The `residuals` are the difference between original image projections and estimated projections (the square of 2-norm of this vector corresponds to the error we computed in Q3.2):

```
residuals = numpy.concatenate([(p1-p1_hat).reshape([-1]),
                 (p2-p2_hat).reshape([-1])]).
```

Use this error function and Scipy's nonlinear least square optimizer `leastsq` write a function to optimize for the best extrinsic matrix and 3D points using the inlier correspondences from `some_corresp_noisy.npz` and the RANSAC estimate of the extrinsics and 3D points as an initialization

```
[M2, w] = bundleAdjustment(K1, M1, p1, K2, M2_init, p2, w_init).
```

In your write-up: include an image of the original 3D points and the optimized points as well as the reprojection error with your initial $\mathbf{M}_2$ and $\mathbf{w}$, and with the optimized matrices.