

16-720A — Spring 2021 — Homework 6

Jen-Hung Ho
jenhungh@andrew.cmu.edu

May 7, 2021

1 Calibrated photometric stereo

(a) **Understanding $n\text{-dot-}l$ lighting.**

The $n\text{-dot-}l$ lighting model:

$$L = \frac{\rho_d}{\pi} A(\hat{\mathbf{n}} \cdot \hat{\mathbf{l}}) = \frac{\rho_d}{\pi} A \cos(\theta) \quad (1)$$

where θ is the angle between the lighting direction and the normal vector.

Since the normal vector $\hat{\mathbf{n}}$ and the lighting direction $\hat{\mathbf{l}}$ are both unit vector, their dot product is $\cos(\theta)$. The projected area is $A \cos(\theta)$, where A is the original area. The viewing direction doesn't matter because the brightness of the whole surface is equal.

(b) **Rendering $n \cdot l$ lighting.**




Figure 1: The appearance of the sphere with incoming lighting direction $(1, 1, 1)/\sqrt{3}$




Figure 2: The appearance of the sphere with incoming lighting direction $(1, -1, 1)/\sqrt{3}$




Figure 3: The appearance of the sphere with incoming lighting direction $(-1, -1, 1)/\sqrt{3}$

(d) **Initials.**

The rank of \mathbf{I} should be 3 because the model is in 3D coordinates, and thus the minimum lighting directions for reconstruction is 3. The size of \mathbf{I} , \mathbf{L} , and \mathbf{B} are $(3, P)$, $(3, 3)$, and $(3, P)$. However, the singular values of $\mathbf{I} = \begin{bmatrix} 79.36348099 & 13.16260675 & 9.22148403 & 2.414729 & 1.61659626 \\ 1.26289066 & 0.89368302 & & & \end{bmatrix}$, which does NOT agree with the rank-3 requirement. This is probably because of the noises or blur of the real world images. Hence, we need more lighting directions (7 lighting directions for this model) to reconstruct the surface.

(e) **Estimating pseudonormals.**

In general, a Least-Squares Problem is formed as $\mathbf{Ax} = \mathbf{y}$, and the estimation is the pseudo inverse $\mathbf{x} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$. In this problem, we treat the relation $\mathbf{L}^T \mathbf{B} = \mathbf{I}$ as a Least-Squares and construct \mathbf{L}^T as matrix \mathbf{A} and \mathbf{I} as vector \mathbf{y} . Therefore, the estimation of the pseudonormals is $\mathbf{B} = (\mathbf{L} \mathbf{L}^T)^{-1} \mathbf{L} \mathbf{I}$.

(f) **Albedos and normals.**




Figure 4: Visualization of the Albedos Image (gray colormap)




Figure 5: Visualization of the Normals Image (rainbow colormap)

From Figure 4, we can discover that the albedos image couldn't show the hair and performs poorly around the ears and nose. Since we only consider 7 lighting directions and the $n \cdot l$ algorithm couldn't handle shadows, we might lose some information if the area is under shadows.

(g) Normals and depth.

The equation of a surface is

$$n_1 * x + n_2 * y + n_3 * z + D = n_1 * x + n_2 * y + n_3 * f(x, y) + D = 0 \quad (2)$$

where the normal vector at point (x, y, z) is $\mathbf{n} = (n_1, n_2, n_3)$

Take partial derivatives of f at (x, y) :

$$f(x, y) = -\frac{D}{n_3} - \frac{n_1}{n_3}x - \frac{n_2}{n_3}y \quad (3)$$

$$f_x = \frac{\partial f(x, y)}{\partial x} = -\frac{n_1}{n_3} \quad (4)$$

$$f_y = \frac{\partial f(x, y)}{\partial y} = -\frac{n_2}{n_3} \quad (5)$$

Therefore, we know that \mathbf{n} is related to the partial derivatives of f at (x, y) .

(h) Understanding integrability of gradients.

The gradients of the 2D and discrete function g are:

$$g_x = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \quad (6)$$

$$g_y = \begin{bmatrix} 4 & 4 & 4 & 4 \\ 4 & 4 & 4 & 4 \\ 4 & 4 & 4 & 4 \\ 4 & 4 & 4 & 4 \end{bmatrix} \quad (7)$$

Reconstruct function g using the two procedures, we can get the same reconstructed g :

$$g_x = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 13 & 14 & 15 & 16 \end{bmatrix} \quad (8)$$

If we modify the first row or column into a non-integrable function, then the estimated functions g in each procedure are not the same, which makes g_x and g_y non-integrable.

The gradients estimated in the way of (g) might be non-integrable because the partial derivatives with respect to x or y might not exist in some special functions.

(i) **Shape estimation.**




Figure 6: Visualization of the Reconstructed Surface (coolwarm colormap)

2 Uncalibrated photometric stereo

(a) **Uncalibrated normal estimation.**

We could compute the best rank-3 approximation of matrix \mathbf{I}

$$\mathbf{I} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = (\mathbf{U}\mathbf{\Sigma}^{1/2})(\mathbf{\Sigma}^{1/2}\mathbf{V}^T) = \hat{\mathbf{L}}^T \hat{\mathbf{B}} \quad (9)$$

Set all singular values except the top 3 from $\mathbf{\Sigma}$ to 0 and extract the first 3 rows of $\hat{\mathbf{L}}$ and $\hat{\mathbf{B}}$. The shape of $\hat{\mathbf{L}}$ and $\hat{\mathbf{B}}$ are $(3, 3)$ and $(3, P)$.

(b) **Calculation and visualization.**



Figure 7: Visualization of the Albedos Image (gray colormap)




Figure 8: Visualization of the Normals Image (rainbow colormap)

(c) **Comparing to ground truth lighting.**

The ground truth lighting direction $\mathbf{L}_0 =$

$$\begin{bmatrix} -0.1418 & 0.1215 & -0.069 & 0.067 & -0.1627 & 0. & 0.1478 \\ -0.1804 & -0.2026 & -0.0345 & -0.0402 & 0.122 & 0.1194 & 0.1209 \\ -0.9267 & -0.9717 & -0.838 & -0.9772 & -0.979 & -0.9648 & -0.9713 \end{bmatrix} \quad (10)$$

The estimated lighting direction $\hat{\mathbf{L}} =$

$$\begin{bmatrix} -2.9927 & -3.87 & -2.408 & -3.745 & -3.5914 & -3.3867 & -3.3525 \\ 0.9478 & -2.3171 & 0.4991 & -0.626 & 2.3257 & 0.4661 & -0.7927 \\ 1.8793 & 1.0146 & 0.4294 & -0.0173 & -0.3108 & -0.9127 & -1.883 \end{bmatrix} \quad (11)$$

The ground truth and estimated lighting directions are NOT similar. There are multiple factorization methods to choose. For example, set \mathbf{L}^T as $\mathbf{U}\Sigma$ and \mathbf{B} as \mathbf{V}^T . As long as the relation $\mathbf{I} = \mathbf{L}^T \mathbf{B}$ doesn't change, the output images are the same.

(d) **Reconstructing the shape, attempt 1.**




Figure 9: Visualization of the Reconstructed Surface (coolwarm colormap)

According to Figure 9, the reconstructed surface does NOT look like a face.


(e) **Reconstructing the shape, attempt 2.**

Figure 10: Visualization of the Reconstructed Surface (coolwarm colormap)

According to Figure 10, the reconstructed surface looks similar to the one output by calibrated photometric stereo.

(f) **Why low relief?**




Figure 11: Vary the parameters μ in the bas-relief transformation ($\mu = -2, 1, 10$)

From Figure 11, we could guess that parameter μ affects the gradients of the reconstructed surfaces since the range of the z axis increases with the values of μ .




Figure 12: Vary the parameters ν in the bas-relief transformation ($\nu = -2, 1, 10$)

From Figure 12, we could guess that parameter ν affects the gradients of the reconstructed surfaces since the range of the z axis increases with the values of ν . Also, ν affects the direction of the reconstructed surface. When ν is negative, the surface is upside down.




Figure 13: Vary the parameters λ in the bas-relief transformation ($\lambda = 0.5, 1, 10$)

From Figure 12, we could guess that parameter λ affects the flatness of the reconstructed surfaces since the surface is flatter when λ is large.

(g) **Flattest surface possible.**

If we want to make the estimated surface as flat as possible, I would increase λ (increase flatness) and set μ and ν to 0 (minimize the gradients).

(h) **More measurements.**

Acquiring more pictures from more lighting directions will help resolve the ambiguity since we have more information about the surface.

Code Appendix

2.1 q1.py

```

1  # ##### #
2  # 16720: Computer Vision Homework 6
3  # Carnegie Mellon University
4  # April 22, 2021
5  # ##### #
6
7  import numpy as np
8  from matplotlib import pyplot as plt
9  import skimage.io
10 from skimage.color import rgb2xyz
11 from utils import integrateFrankot, plotSurface
12
13 def renderNdotLSphere(center, rad, light, pxSize, res):
14
15     """
16     Question 1 (b)
17
18     Render a hemispherical bowl with a given center and radius. Assume that
19     the hollow end of the bowl faces in the positive z direction, and the
20     camera looks towards the hollow end in the negative z direction. The
21     camera's sensor axes are aligned with the x- and y-axes.
22
23     Parameters
24     -----
25     center : numpy.ndarray
26         The center of the hemispherical bowl in an array of size (3,)
27
28     rad : float
29         The radius of the bowl
30
31     light : numpy.ndarray
32         The direction of incoming light
33
34     pxSize : float
35         Pixel size
36
37     res : numpy.ndarray
38         The resolution of the camera frame
39
40     Returns
41     -----
42     image : numpy.ndarray
43         The rendered image of the hemispherical bowl
44     """

```

```

45     [X, Y] = np.meshgrid(np.arange(res[0]), np.arange(res[1]))
46     X = (X - res[0]/2) * pxSize * 1.e-4
47     Y = (Y - res[1]/2) * pxSize * 1.e-4
48     Z = np.sqrt(rad**2 + 0j - X**2 - Y**2)
49     Z = np.real(Z)
50
51
52     image = None
53     # Your code here
54     # Build the normal vector
55     n_vector = np.concatenate((X[:, :, np.newaxis], Y[:, :, np.newaxis],
56                               Z[:, :, np.newaxis]), axis = 2)
57     N = n_vector.reshape(res[0]*res[1], -1)
58     # Normalization
59     norm = np.linalg.norm(N, ord = 2, axis = 1)
60     N = N / norm.reshape(-1, 1)
61
62     # Implement NdotL Algorithm
63     L = light
64     image = (N @ L).reshape(res[1], res[0])
65     image[np.real(Z) == 0] = 0
66
67     return image
68
69
70 def loadData(path = "../data/"):
71
72     """
73     Question 1 (c)
74
75     Load data from the path given. The images are stored as input_n.tif
76     for n = {1...7}. The source lighting directions are stored in
77     sources.npy.
78
79     Parameters
80     -----
81     path: str
82         Path of the data directory
83
84     Returns
85     -----
86     I : numpy.ndarray
87         The 7 x P matrix of vectorized images
88
89     L : numpy.ndarray
90         The 3 x 7 matrix of lighting directions
91

```

```

92     s: tuple
93         Image shape
94
95     """
96
97     I = None
98     L = None
99     s = None
100
101     # Your code here
102     # Load the image and Compute I and L
103     num_img = 7
104     for i in range(1, num_img+1):
105         # Load the image and check the datatype
106         input_img_rgb = skimage.io.imread(path + f'input_{i}.tif')
107         input_img_rgb = input_img_rgb.astype(np.uint16)
108
109         # Convert the RGB images into the XYZ color space
110         input_img_xyz = rgb2xyz(input_img_rgb)
111
112         # Compute s
113         h, w, _ = input_img_rgb.shape
114         s = (h, w)
115
116         # Compute L
117         if I is None:
118             I = np.zeros((7, h*w))
119             I[i-1, :] = input_img_xyz[:, :, 1].reshape(1, h*w)
120
121         # Compute I
122         L = np.load(path + 'sources.npy')
123         L = L.T
124
125     return I, L, s
126
127
128 def estimatePseudonormalsCalibrated(I, L):
129
130     """
131     Question 1 (e)
132
133     In calibrated photometric stereo, estimate pseudonormals from the
134     light direction and image matrices
135
136     Parameters
137     -----
138     I : numpy.ndarray
139         The 7 x P array of vectorized images

```



```

140
141     L : numpy.ndarray
142         The 3 x 7 array of lighting directions
143
144     Returns
145     -----
146     B : numpy.ndarray
147         The 3 x P matrix of pseudonormals
148     """
149
150     B = None
151     # Your code here
152     # Least Square Problem :  $Ax = y$ 
153     #  $L.T @ B = I$ 
154     A = L.T
155     y = I
156     # Pseudo-inverse :  $x = \text{inv}(A.T @ A) @ A.T @ y$ 
157     B = np.linalg.inv(A.T @ A) @ A.T @ y
158
159     return B
160
161
162 def estimateAlbedosNormals(B):
163
164     '''
165     Question 1 (e)
166
167     From the estimated pseudonormals, estimate the albedos and normals
168
169     Parameters
170     -----
171     B : numpy.ndarray
172         The 3 x P matrix of estimated pseudonormals
173
174     Returns
175     -----
176     albedos : numpy.ndarray
177         The vector of albedos
178
179     normals : numpy.ndarray
180         The 3 x P matrix of normals
181     '''
182
183     albedos = None
184     normals = None
185     # Your code here
186     # albedos = the magnitudes of the pseudonormals

```

```

187     albedos = np.linalg.norm(B, ord = 2, axis = 0)
188
189     # normals = normalized normal vectors
190     normals = B / albedos
191
192     return albedos, normals
193
194
195 def displayAlbedosNormals(albedos, normals, s):
196
197     """
198     Question 1 (f, g)
199
200     From the estimated pseudonormals, display the albedo and normal maps
201
202     Please make sure to use the `coolwarm` colormap for the albedo image
203     and the `rainbow` colormap for the normals.
204
205     Parameters
206     -----
207     albedos : numpy.ndarray
208         The vector of albedos
209
210     normals : numpy.ndarray
211         The 3 x P matrix of normals
212
213     s : tuple
214         Image shape
215
216     Returns
217     -----
218     albedoIm : numpy.ndarray
219         Albedo image of shape s
220
221     normalIm : numpy.ndarray
222         Normals reshaped as an s x 3 image
223
224     """
225
226     albedoIm = None
227     normalIm = None
228     # Your code here
229     # Reshape albedos
230     albedoIm = albedos.reshape(s)
231
232     # Rescale and Reshape normals
233     normals = (normals + 1) / 2
234     normalIm = normals.T.reshape(s[0], s[1], 3)

```

```

235
236     return albedoIm, normalIm
237
238
239 def estimateShape(normals, s):
240
241     """
242     Question 1 (j)
243
244     Integrate the estimated normals to get an estimate of the depth map
245     of the surface.
246
247     Parameters
248     -----
249     normals : numpy.ndarray
250         The 3 x P matrix of normals
251
252     s : tuple
253         Image shape
254
255     Returns
256     -----
257     surface: numpy.ndarray
258         The image, of size s, of estimated depths at each point
259
260     """
261
262     surface = None
263     # Your code here
264     # Rescale and Reshape normals
265     # normals = (normals + 1) / 2
266     normalIm = normals.T.reshape(s[0], s[1], 3)
267
268     # Compute the partial derivatives
269     n1, n2, n3 = normalIm[:, :, 0], normalIm[:, :, 1], normalIm[:, :, 2]
270     df_dx = -n1 / n3
271     df_dy = -n2 / n3
272
273     # Estimate the actual surface
274     surface = integrateFrankot(df_dx, df_dy)
275
276     return surface
277
278
279 if __name__ == '__main__':
280     # Part 1(b)
281     radius = 0.75 # cm
282     center = np.asarray([0, 0, 0]) # cm

```

```

283     pxSize = 7 # um
284     res = (3840, 2160)
285
286     light = np.asarray([1, 1, 1])/np.sqrt(3)
287     image = renderNDotLSphere(center, radius, light, pxSize, res)
288     plt.figure()
289     plt.imshow(image, cmap = 'gray', vmin = 0, vmax = 1)
290     plt.show()
291     plt.imsave('1b-a.png', image, cmap = 'gray', vmin = 0, vmax = 1)
292
293     light = np.asarray([1, -1, 1])/np.sqrt(3)
294     image = renderNDotLSphere(center, radius, light, pxSize, res)
295     plt.figure()
296     plt.imshow(image, cmap = 'gray', vmin = 0, vmax = 1)
297     plt.show()
298     plt.imsave('1b-b.png', image, cmap = 'gray', vmin = 0, vmax= 1)
299
300     light = np.asarray([-1, -1, 1])/np.sqrt(3)
301     image = renderNDotLSphere(center, radius, light, pxSize, res)
302     plt.figure()
303     plt.imshow(image, cmap = 'gray', vmin = 0, vmax = 1)
304     plt.show()
305     plt.imsave('1b-c.png', image, cmap = 'gray', vmin = 0, vmax = 1)
306
307     # Part 1(c)
308     I, L, s = loadData('../data/')
309
310     # Part 1(d)
311     # Singular Value Decomposition of I
312     u, sin, vh = np.linalg.svd(I, full_matrices = False)
313     print(f"Singular Value of I = {sin}")
314
315     # Part 1(e)
316     B = estimatePseudonormalsCalibrated(I, L)
317
318     # # Part 1(f)
319     albedos, normals = estimateAlbedosNormals(B)
320     albedoIm, normalIm = displayAlbedosNormals(albedos, normals, s)
321     plt.imsave('1f-a.png', albedoIm, cmap = 'gray')
322     plt.imsave('1f-b.png', normalIm, cmap = 'rainbow')
323
324     # # Part 1(i)
325     surface = estimateShape(normals, s)
326     plotSurface(surface, suffix = '1i')

```

2.2 q2.py

```

1  # #####
2  # 16720: Computer Vision Homework 6
3  # Carnegie Mellon University
4  # April 22, 2021
5  # #####
6
7  import numpy as np
8  import matplotlib.pyplot as plt
9  from q1 import loadData, estimateAlbedosNormals, displayAlbedosNormals,
10                                     estimateShape
11 from q1 import estimateShape
12 from utils import enforceIntegrability, plotSurface
13
14 def estimatePseudonormalsUncalibrated(I):
15
16     """
17     Question 2 (b)
18
19     Estimate pseudonormals without the help of light source directions.
20
21     Parameters
22     -----
23     I : numpy.ndarray
24         The 7 x P matrix of loaded images
25
26     Returns
27     -----
28     B : numpy.ndarray
29         The 3 x P matrix of pseudonormals
30
31     L : numpy.ndarray
32         The 3 x 7 array of lighting directions
33
34     """
35
36     B = None
37     L = None
38     # Your code here
39     # Singular Value Decomposition of I
40     u, sin, vh = np.linalg.svd(I, full_matrices = False)
41
42     # Estimate the best rank-3 approximation
43     # Set all other singular values to 0
44     sin[3:] = 0
45     s_sqrt = np.sqrt(np.diag(sin))
46

```

```

47     # Estimate B and L :  $L.T @ B = I$ 
48     B = (s_sqrt @ vh)[0:3, :]
49     L = (u @ s_sqrt).T[0:3, :]
50
51     return B, L
52
53 def plotBasRelief(B, mu, nu, lam, suffix = ''):
54
55     """
56     Question 2 (f)
57
58     Make a 3D plot of of a bas-relief transformation with the given parameters.
59
60     Parameters
61     -----
62     B : numpy.ndarray
63         The 3 x P matrix of pseudonormals
64
65     mu : float
66         bas-relief parameter
67
68     nu : float
69         bas-relief parameter
70
71     lambda : float
72         bas-relief parameter
73
74     Returns
75     -----
76     None
77
78     """
79
80     # Your code here
81     # Form the G matrix
82     G = np.array([[1, 0, 0], [0, 1, 0], [mu, nu, lam]])
83
84     # Generalized bas-relief ambiguity
85     new_B = np.linalg.inv(G).T @ B
86
87     # Visulaize the reconstructed 3D depth map
88     integrable_normals = enforceIntegrability(new_B, s)
89     surface = estimateShape(integrable_normals, s)
90     plotSurface(surface, suffix = suffix)
91
92
93 if __name__ == "__main__":

```

```

94
95     # Part 2 (b)
96     # Estimate B_hat and L_hat
97     I, L0, s = loadData()
98     B_hat, L_hat = estimatePseudonormalsUncalibrated(I)
99     # Visualize the estimated albedos and normals
100    albedos, normals = estimateAlbedosNormals(B_hat)
101    albedoIm, normalIm = displayAlbedosNormals(albedos, normals, s)
102    plt.imsave('2b-a.png', albedoIm, cmap = 'gray')
103    plt.imsave('2b-b.png', normalIm, cmap = 'rainbow')
104
105    # Part 2 (c)
106    # Compare L0 and L_hat
107    print(f"Ground truth lighting L0 =\n{L0}")
108    print(f"Estimated lighting L_hat =\n{L_hat}")
109
110    # Part 2 (d)
111    # Visualize the reconstructed 3D depth map
112    surface = estimateShape(normals, s)
113    plotSurface(surface, suffix = '2d')
114
115    # Part 2 (e)
116    # Transform non-integrable pseudonormals into integrable pseudonormals
117    integrable_normals = enforceIntegrability(B_hat, s)
118    # Visualize the reconstructed 3D depth map
119    surface = estimateShape(integrable_normals, s)
120    plotSurface(surface, suffix = '2e')
121
122    # Part 2 (f)
123    # Visualize the corresponding surfaces
124    vary_parameters = [-2, 1, 10]
125    for i in vary_parameters:
126        plotBasRelief(B_hat, i, 0, 1, suffix = f'_mu_{i}')
127        plotBasRelief(B_hat, 0, i, 1, suffix = f'_nu_{i}')
128        if i > 0:
129            plotBasRelief(B_hat, 0, 0, i, suffix = f'_lam_{i}')
```

2.3 utils.py

```

1  # #####
2  # 16720: Computer Vision Homework 6
3  # Carnegie Mellon University
4  # April 22, 2021
5  # #####
6
7  import numpy as np
8  import warnings
9  from scipy.ndimage import gaussian_filter
10 from matplotlib import pyplot as plt
11 from matplotlib import cm
12
13 def integrateFrankot(zx, zy, pad = 512):
14
15     """
16     Question 1 (j)
17
18     Implement the Frankot-Chellappa algorithm for enforcing integrability
19     and normal integration
20
21     Parameters
22     -----
23     zx : numpy.ndarray
24         The image of derivatives of the depth along the x image dimension
25
26     zy : tuple
27         The image of derivatives of the depth along the y image dimension
28
29     pad : float
30         The size of the full FFT used for the reconstruction
31
32     Returns
33     -----
34     z: numpy.ndarray
35         The image, of the same size as the derivatives, of estimated depths
36         at each point
37
38     """
39
40     # Raise error if the shapes of the gradients don't match
41     if not zx.shape == zy.shape:
42         raise ValueError('Sizes of both gradients must match!')
43
44     # Pad the array FFT with a size we specify
45     h, w = 512, 512
46

```



```

47     # Fourier transform of gradients for projection
48     Zx = np.fft.fftshift(np.fft.fft2(zx, (h, w)))
49     Zy = np.fft.fftshift(np.fft.fft2(zy, (h, w)))
50     j = 1j
51
52     # Frequency grid
53     [wx, wy] = np.meshgrid(np.linspace(-np.pi, np.pi, w),
54                             np.linspace(-np.pi, np.pi, h))
55     absFreq = wx**2 + wy**2
56
57     # Perform the actual projection
58     with warnings.catch_warnings():
59         warnings.simplefilter('ignore')
60         z = (-j*wx*Zx-j*wy*Zy)/absFreq
61
62     # Set (undefined) mean value of the surface depth to 0
63     z[0, 0] = 0.
64     z = np.fft.ifftshift(z)
65
66     # Invert the Fourier transform for the depth
67     z = np.real(np.fft.ifft2(z))
68     z = z[:zx.shape[0], :zx.shape[1]]
69
70     return z
71
72
73 def enforceIntegrability(N, s, sig = 3):
74
75     """
76     Question 2 (e)
77
78     Find a transform Q that makes the normals integrable and transform them
79     by it
80
81     Parameters
82     -----
83     N : numpy.ndarray
84         The 3 x P matrix of (possibly) non-integrable normals
85
86     s : tuple
87         Image shape
88
89     Returns
90     -----
91     Nt : numpy.ndarray
92         The 3 x P matrix of transformed, integrable normals
93     """

```

```

95     N1 = N[0, :].reshape(s)
96     N2 = N[1, :].reshape(s)
97     N3 = N[2, :].reshape(s)
98
99     N1y, N1x = np.gradient(gaussian_filter(N1, sig), edge_order = 2)
100    N2y, N2x = np.gradient(gaussian_filter(N2, sig), edge_order = 2)
101    N3y, N3x = np.gradient(gaussian_filter(N3, sig), edge_order = 2)
102
103    A1 = N1*N2x-N2*N1x
104    A2 = N1*N3x-N3*N1x
105    A3 = N2*N3x-N3*N2x
106    A4 = N2*N1y-N1*N2y
107    A5 = N3*N1y-N1*N3y
108    A6 = N3*N2y-N2*N3y
109
110    A = np.hstack((A1.reshape(-1, 1),
111                  A2.reshape(-1, 1),
112                  A3.reshape(-1, 1),
113                  A4.reshape(-1, 1),
114                  A5.reshape(-1, 1),
115                  A6.reshape(-1, 1)))
116
117    AtA = A.T.dot(A)
118    W, V = np.linalg.eig(AtA)
119    h = V[:, np.argmin(np.abs(W))]
120
121    delta = np.asarray([[ -h[2],  h[5], 1],
122                       [ h[1], -h[4], 0],
123                       [-h[0],  h[3], 0]])
124    Nt = np.linalg.inv(delta).dot(N)
125
126    return Nt
127
128 def plotSurface(surface, suffix=''):
129
130     """
131     Plot the depth map as a surface
132
133     Parameters
134     -----
135     surface : numpy.ndarray
136         The depth map to be plotted
137
138     suffix: str
139         suffix for save file
140
141     Returns
142     -----

```

```
143         None
144
145     """
146     x, y = np.meshgrid(np.arange(surface.shape[1]),
147                        np.arange(surface.shape[0]))
148     fig = plt.figure()
149     ax = fig.gca(projection='3d')
150     surf = ax.plot_surface(x, y, -surface, cmap = cm.coolwarm,
151                           linewidth = 0, antialiased = False)
152     ax.view_init(elev = 60., azimuth = 75.)
153     plt.savefig(f'faceCalibrated{suffix}.png')
154     plt.show()
```